

Advanced Machine Learning with scikit-learn Part II/II

Instructor

- [Andreas Mueller @amuellerm1](#) - Columbia University; [Book: Introduction to Machine Learning with Python](#)

This repository will contain the teaching material and other info associated with the workshop "Advanced Machine Learning with scikit-learn Part II/II".

Please download the Large Movie Review dataset from <http://ai.stanford.edu/~amaas/data/sentiment/> before coming to the workshop!

About the workshop

Scikit-learn is a machine learning library in Python, that has become a valuable tool for many data science practitioners. This training will cover some advanced topics in using scikit-learn, such as how to perform out-of-core learning with scikit-learn and how to speed up parameter search. We'll also cover how to build your own models or feature extraction methods that are compatible with scikit-learn, which is important for feature extraction in many domains. We will see how we can customize scikit-learn even further, using custom methods for cross-validation or model evaluation.

Prerequisites

This workshop assumes familiarity with Jupyter notebooks and basics of pandas, matplotlib and numpy. It also assumes experience using scikit-learn and familiarity with the API.

Content

- Out-of-core learning with scikit-learn
- Feature Hashing
- Building custom estimators and extending scikit-learn
- Advanced features in scikit-learn
- Efficient parameter tuning
- TBA

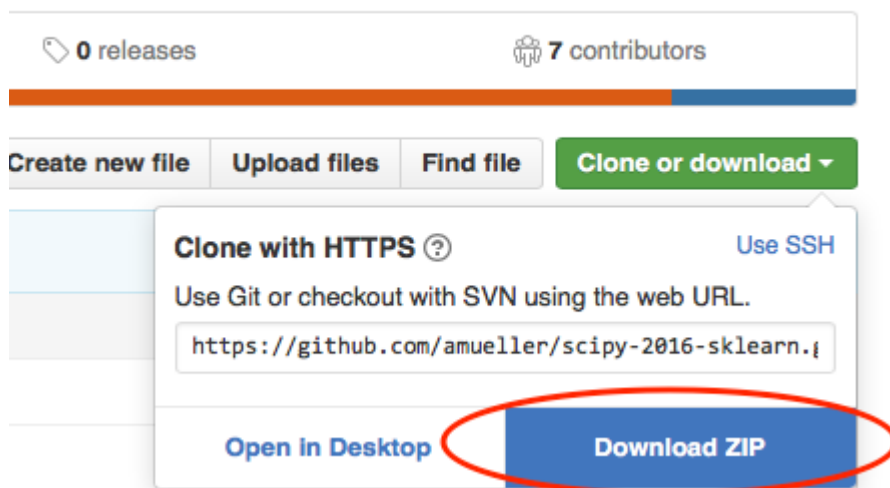
Obtaining the Tutorial Material

If you are familiar with git, it is most convenient if you clone the GitHub repository. This is highly encouraged as it allows you to easily synchronize any changes to the material.

```
git clone https://github.com/amueller/ml-workshop-4-of-4
```

If you are not familiar with git, you can download the repository as a .zip file by heading over to the GitHub repository (<https://github.com/amueller/ml-workshop-4-of-4>) in your browser and click the green "Download"

button in the upper right.



Please note that I may add and improve the material until shortly before the tutorial session, and we recommend you to update your copy of the materials one day before the tutorials. If you have an GitHub account and forked/cloned the repository via GitHub, you can sync your existing fork with via the following commands:

```
git pull origin master
```

Installation Notes

This tutorial will require recent installations of

- [NumPy](#)
- [SciPy](#)
- [matplotlib](#)
- [pillow](#)
- [pandas](#)
- [scikit-learn](#) ($\geq 0.18.1$)
- [IPython](#)
- [Jupyter Notebook](#)
- [mlxtend](#)
- [imbalance-learn](#)

The last one is important, you should be able to type:

```
jupyter notebook
```

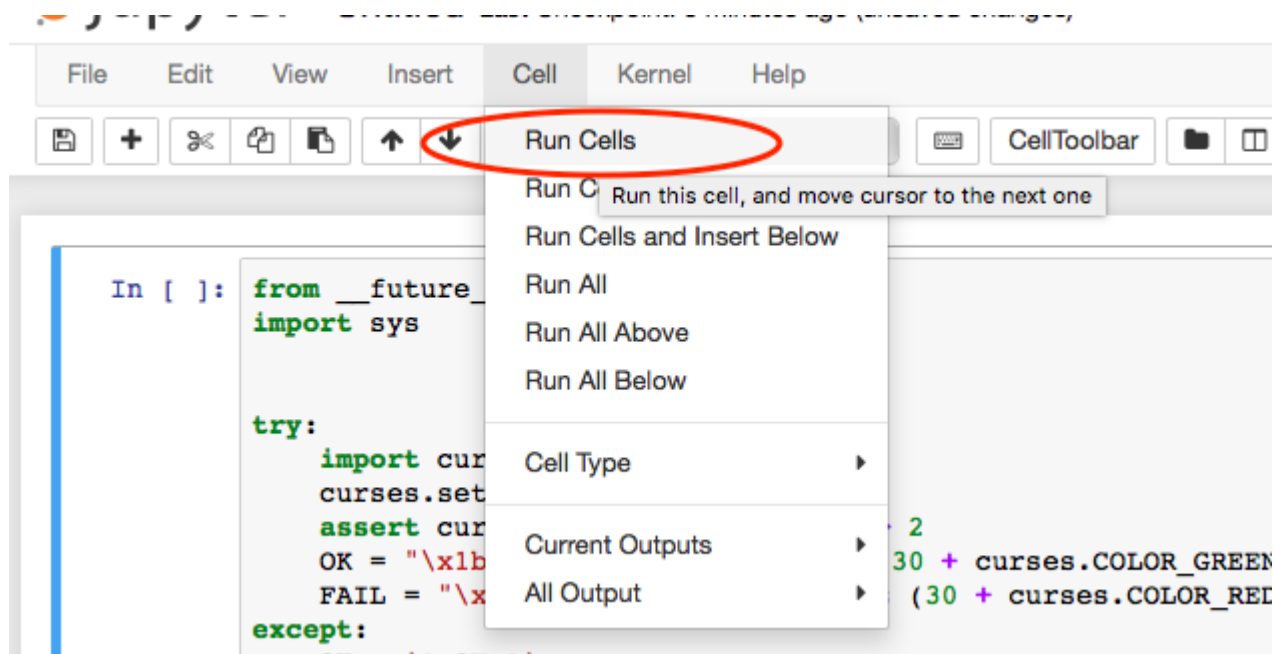
in your terminal window and see the notebook panel load in your web browser. Try opening and running a notebook from the material to see check that it works.

For users who do not yet have these packages installed, a relatively painless way to install all the requirements is to use a Python distribution such as [Anaconda](#), which includes the most relevant Python packages for science, math, engineering, and data analysis; Anaconda can be downloaded and installed for free including commercial use and redistribution. The code examples in this tutorial should be compatible to Python 2.7, Python 3.4 and later. However, it's recommended to use a recent Python version (like 3.5 or 3.6).

After obtaining the material, we **strongly recommend** you to open and execute a Jupyter Notebook [jupyter notebook check_env.ipynb](#) that is located at the top level of this repository. Inside the repository, you can open the notebook by executing

```
jupyter notebook check_env.ipynb
```

inside this repository. Inside the Notebook, you can run the code cell by clicking on the "Run Cells" button as illustrated in the figure below:



Finally, if your environment satisfies the requirements for the tutorials, the executed code cell will produce an output message as shown below:

```
Using python in /Users/Sebastian/miniconda3
3.5.1 |Continuum Analytics, Inc.| (default, Jun 15 2016, 16:14:02)
[GCC 4.2.1 Compatible Apple LLVM 4.2 (clang-425.0.28)]

[ OK ] IPython version 4.2.0
[ OK ] numpy version 1.11.0
[ OK ] watermark version 1.3.1
[ OK ] matplotlib version 1.5.1
[ OK ] scipy version 0.17.1
[ OK ] yaml version 3.11
[ OK ] PIL version 1.1.7
[ OK ] sklearn version 0.17.1
[ OK ] pydot version 1.2.2
```